

BRACT's

Vishwakarma Institute of Technology, Pune

Practical Implementation Sheet

Department: CSE-AI Semester: 5 Academic Year: 2026-27

Division: E Batch: 1 Practical No: 5

Roll no: 8 PRN No: 12412097 Name of Student: Sahil Balasaheb Patil

Subject Name : Deep Learning

Problem Statement :

Implement and compare RNN, LSTM, GRU Model for sequence classification and analyze their performance using appropriate evaluation metrics

Algorithm :

1. Import TensorFlow and required libraries.
2. Load the MNIST handwritten digit dataset.
3. Split the dataset into training and testing sets.
4. Normalize pixel values from 0–255 to 0–1.
5. Build a neural network model using Keras.
6. Define the input, hidden, and output layers.
7. Compile the model using an optimizer, loss function, and accuracy metric.
8. Train the model using the training images and labels.
9. Evaluate the trained model on the test dataset.
10. Test the model on sample images and display the predicted digits.

Code and Output

-

Implement and compare RNN, LSTM, and GRU models for binary sequence classification using the IMDb 50k movie review dataset and evaluate them using accuracy, precision, recall, F1-score, confusion matrices, training time, and parameter count.

```
[ ]  
  
# Install the libraries required for dataset downloading,  
# data manipulation, deep learning, evaluation, and visualization.  
  
!pip install -q kagglehub tensorflow scikit-learn seaborn pandas matplotlib  
  
[ ]  
  
# Import operating-system utilities for handling dataset file paths.  
import os  
  
# Import regular expressions for text cleaning.  
import re  
  
# Import time to measure model training duration.  
import time  
  
# Import random to control random operations.  
import random  
  
# Import NumPy for numerical operations.  
import numpy as np  
  
# Import Pandas for creating and manipulating dataframes.  
import pandas as pd  
  
# Import functions for splitting the dataset.  
from sklearn.model_selection import train_test_split  
  
# Import evaluation metrics for binary classification.  
from sklearn.metrics import (  
    accuracy_score,  
    precision_score,  
    recall_score,  
    f1_score,  
    confusion_matrix,  
    classification_report,  
    roc_curve,  
    auc,  
    roc_auc_score,  
    precision_recall_curve,  
    average_precision_score  
)  
  
# Import the Keras tokenizer for converting text into integer sequences.  
from tensorflow.keras.preprocessing.text import Tokenizer  
  
# Import padding utility to make all sequences equal in length.  
from tensorflow.keras.preprocessing.sequence import pad_sequences  
  
# Import the Sequential model API.  
from tensorflow.keras.models import Sequential  
  
# Import layers required for the three models.  
from tensorflow.keras.layers import (  
    Input,  
    Embedding,  
    SimpleRNN,  
    LSTM,  
    GRU,  
    Dense,  
    Dropout  
)  
  
# Import EarlyStopping to stop training when validation performance stops improving.  
from tensorflow.keras.callbacks import EarlyStopping
```

+ Code

+ Text

Start coding or generate with AI.

SEED = 42

```
np.random.seed(SEED)  
random.seed(SEED)  
tf.random.set_seed(SEED)
```

```
sns.set_theme(style="whitegrid")
```

```
# Dataset: IMDB Movie Review Dataset(kaggle)  
import kagglehub  
  
path = kagglehub.dataset_download(  
    "lakshmi25npathi/imdb-dataset-of-50k-movie-reviews"  
)  
  
print("Dataset downloaded to:")  
print(path)
```

```
print("\nFiles:")
print(os.listdir(path))
```

Using Colab cache for faster access to the 'imdb-dataset-of-50k-movie-reviews' dataset.
Dataset downloaded to:
/kaggle/input/imdb-dataset-of-50k-movie-reviews

Files:
['IMDB Dataset.csv']

```
# Load Dataset:
file_path = os.path.join(
    path,
    "IMDB Dataset.csv"
)

df = pd.read_csv(file_path)

print("Dataset shape:", df.shape)

df.head()
```

Dataset shape: (50000, 2)

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Matte's "Love in the Time of Money" is...	positive

```
# Display basic information about the dataset,
# including column names, data types, and non-null values.
```

```
# including column names, data types, and non-null values.
print("Dataset Information")
print("=" * 60)
df.info()

# Display the number of missing values in each column.
print("\nMissing Values:")
print(df.isnull().sum())

# Count duplicate reviews.
# Duplicate reviews may cause data leakage if the same review
# appears in both training and testing sets.
print("\nDuplicate Reviews:")
print(df["review"].duplicated().sum())

# Display the number of positive and negative reviews.
# This helps us understand the class distribution.
print("\nSentiment Distribution:")
print(df["sentiment"].value_counts())
```

*** Dataset Information


```
=====
Number of reviews: 50000
Number of columns: 2
```

```
Columns:
['review', 'sentiment']
```

```
Missing values:
review      0
sentiment   0
dtype: int64
```

```
Duplicate reviews:
418
```

```
Sentiment distribution:
sentiment
-----
positive  35000
```



```
# Define a function to clean individual movie reviews.
def clean_text(text):

    # Convert all characters to lowercase.
    # This reduces the number of unique word representations.
    text = text.lower()

    # Remove HTML line-break tags such as <br />.
    # These tags are present in some IMDb reviews.
    text = re.sub(r"<br\s*/>", " ", text)

    # Remove non-alphabetic characters.
    # This keeps the text simpler for tokenization.
    text = re.sub(r"[^a-zA-Z\s]", " ", text)

    # Replace multiple spaces with a single space.
    text = re.sub(r"\s+", " ", text).strip()

    # Return the cleaned review.
    return text

# Apply the clean_text function to every review in the dataset.
# The cleaned reviews are stored in a new column.
df["clean_review"] = df["review"].apply(clean_text)

# Display the original and cleaned versions of one review.
# This allows us to verify that preprocessing worked correctly.
print("Original Review:")
print(df["review"].iloc[0])

print("\nCleaned Review:")
print(df["clean_review"].iloc[0])
```

```
# Label Encoding
df["label"] = df["sentiment"].map({
    "negative": 0,
    "positive": 1
})

print(df[["sentiment", "label"]].head())
```

```
*** sentiment label
0 positive 1
1 positive 1
2 positive 1
3 negative 0
4 positive 1
```

```
# Analyze review length
df["review_length"] = df["clean_review"].apply(
    lambda x: len(x.split())
)


print(df["review_length"].describe())
```

```
*** count    50000.000000
mean       234.187100
std        173.516096
min         6.000000
25%        128.000000
50%        176.000000
75%        284.000000
max        2494.000000
Name: review_length, dtype: float64
```

```
# Visualization: Review length distribution
plt.figure(figsize=(10, 6))

sns.histplot(
    data=df,

plt.tight_layout()
plt.show()
```



```

# Train/Test Split:
# Training : 80%
# Testing : 10%
# Validation : 10%

X = df["clean_review"].values
y = df["label"].values
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.10,
    random_state=SEED,
    stratify=y
)
X_train, X_val, y_train, y_val = train_test_split(
    X_train,
    y_train,
    test_size=0.111111,
    random_state=SEED,
    stratify=y_train
)
print("Training samples : ", len(X_train))
print("Validation samples : ", len(X_val))
print("Testing samples : ", len(X_test))

... Training samples : 40000
Validation samples : 5000
Testing samples : 5000

#Tokenizer:
# Set the maximum number of words to retain in the vocabulary.
# Only the most frequent words will be assigned individual indices.
VOCAB_SIZE = 20000

# Set the maximum number of tokens for each review sequence.
BATCH_SIZE = 64

# Set the maximum number of training epochs.
EPOCHS = 5

# Set the learning rate for the Adam optimizer.
LEARNING_RATE = 0.

# Create a tokenizer that keeps the most frequent words
# and assigns an OOV token to unknown words.
tokenizer = Tokenizer(
    num_words=VOCAB_SIZE,
    oov_token="<OOV>"
)

# Fit the tokenizer only on the training reviews.
# This prevents information from validation or test data
# from influencing the vocabulary.
tokenizer.fit_on_texts(X_train)

# Display the total number of unique words discovered.
print("Total words in tokenizer:", len(tokenizer.word_index))
# Convert training reviews into sequences of integer word IDs.
X_train_seq = tokenizer.texts_to_sequences(X_train)

# Convert validation reviews into integer sequences.
X_val_seq = tokenizer.texts_to_sequences(X_val)

# Convert testing reviews into integer sequences.
X_test_seq = tokenizer.texts_to_sequences(X_test)

# Display an example of a tokenized review.
print("Example tokenized review:")
print(X_train_seq[0][:20])

# Pad or truncate training sequences to a fixed length.
X_train_seq = pad_sequences(

```

```

# Pad or truncate training sequences to a fixed length.
X_train_pad = pad_sequences(
    X_train_seq,
    maxlen=MAX_LENGTH,
    padding="post",
    truncating="post"
)

# Pad or truncate validation sequences to the same fixed length.
X_val_pad = pad_sequences(
    X_val_seq,
    maxlen=MAX_LENGTH,
    padding="post",
    truncating="post"
)

# Pad or truncate testing sequences to the same fixed length.
X_test_pad = pad_sequences(
    X_test_seq,
    maxlen=MAX_LENGTH,
    padding="post",
    truncating="post"
)

# Display the shapes of the resulting arrays.
# Each review should now contain exactly MAX_LENGTH tokens.
print("Training shape :", X_train_pad.shape)
print("Validation shape :", X_val_pad.shape)
print("Testing shape :", X_test_pad.shape)

```

```

... Training shape : (40000, 250)
Validation shape : (5000, 250)
Testing shape : (5000, 250)

```

```

# Define common hyperparameters:
EMBEDDING_DIM = 128
GET_DUAL_TNN():

# Create a sequential neural-network architecture.
model = Sequential([

    # Define the input shape as a sequence of MAX_LENGTH integer tokens.
    Input(shape=(MAX_LENGTH,)),

    # Convert each integer word ID into a dense embedding vector.
    Embedding(
        input_dim=VOCAB_SIZE,
        output_dim=EMBEDDING_DIM
    ),

    # Process the sequence using a basic Simple RNN layer.
    # The final hidden state is used for classification.
    SimpleRNN(
        RNN_UNITS
    ),

    # Randomly deactivate a fraction of neurons during training
    # to reduce overfitting.
    Dropout(
        DROPOUT_RATE
    ),

    # Produce one probability for binary classification.
    # Sigmoid outputs a value between 0 and 1.
    Dense(
        1,
        activation="sigmoid"
    )
])

# Compile the model using the Adam optimizer.
model.compile(

```

```

▶ # Define a function that creates an LSTM model.
def build_lstm():

    # Create a sequential neural-network architecture.
    model = Sequential([

        # Define the input shape as a sequence of MAX_LENGTH tokens.
        Input(shape=(MAX_LENGTH,)),

        # Convert word IDs into dense embedding vectors.
        Embedding(
            input_dim=VOCAB_SIZE,
            output_dim=EMBEDDING_DIM
        ),

        # Process the sequence using an LSTM layer.
        # LSTM uses memory cells and gates to learn long-term dependencies.
        LSTM(
            RNN_UNITS
        ),

        # Apply dropout to reduce overfitting.
        Dropout(
            DROPOUT_RATE
        ),

        # Produce a probability for the positive class.
        Dense(
            1,
            activation="sigmoid"
        )
    ])

```

```

▶ # Define a function that creates a GRU model.
def build_gru():

    # Create a sequential neural-network architecture.
    model = Sequential([

        # Define the input shape as a sequence of MAX_LENGTH tokens.
        Input(shape=(MAX_LENGTH,)),

        # Convert word IDs into dense embedding vectors.
        Embedding(
            input_dim=VOCAB_SIZE,
            output_dim=EMBEDDING_DIM
        ),

        # Process the sequence using a GRU layer.
        # GRU uses update and reset gates to control information flow.
        GRU(
            RNN_UNITS
        ),

        # Apply dropout to reduce overfitting.
        Dropout(
            DROPOUT_RATE
        ),

        # Produce a probability for binary classification.
        Dense(
            1,
            activation="sigmoid"
        )
    ])

    # Compile the model using the same optimizer and loss
    # as the RNN and LSTM models.

```



```
# Create All Models:
models = {
    "RNN": build_rnn(),
    "LSTM": build_lstm(),
    "GRU": build_gru()
}
```

```
for name, model in models.items():

    print("\n")
    print("-" * 70)
    print(name)
    print("-" * 70)

    model.summary()
```

```
=====
RNN
=====
Model: "sequential"

```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 256, 128)	2,560,000
simple_rnn (SimpleRNN)	(None, 64)	12,352
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 1)	65

```

Total params: 2,572,417 (9.81 MB)
Trainable params: 2,572,417 (9.81 MB)
Non-trainable params: 0 (0.00 B)
```

```
=====
LSTM
=====
Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 256, 128)	2,560,000
lstm (LSTM)	(None, 64)	49,408
dropout_1 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 1)	65

```

Total params: 2,609,473 (9.95 MB)
Trainable params: 2,609,473 (9.95 MB)
Non-trainable params: 0 (0.00 B)
```

```
=====
GRU
=====
Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 256, 128)	2,560,000
gru (GRU)	(None, 64)	37,248
dropout_2 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 1)	65

```

Total params: 2,597,313 (9.91 MB)
Trainable params: 2,597,313 (9.91 MB)
Non-trainable params: 0 (0.00 B)
```

Total params: 2,597,313 (9.91 MB)
Trainable params: 2,597,313 (9.91 MB)
Non-trainable params: 0 (0.00 B)

```
parameter_results = []

for name, model in models.items():

    parameter_results.append({
        "Model": name,
        "Parameters": model.count_params()
    })

parameter_df = pd.DataFrame(
    parameter_results
)

parameter_df
```

```
...      Model  Parameters
0      RNN    2572417
1      LSTM    2609473
2      GRU    2597313
```


```
plt.figure(figsize=(8, 5))
```

```
ax = sns.barplot(
    data=parameter_df,
    x="Model",
    y="Parameters"
```

```
plt.ylabel("Number of Parameters")

for container in ax.containers:
    ax.bar_label(container, fmt="%.0F")

plt.tight_layout()
plt.show()
```



TRAINING RNN

```
Epoch 1/5
625/625 — 77s 121ms/step - accuracy: 0.5021 - loss: 0.6990 - val_accuracy: 0.5088 - val_loss: 0.6930
Epoch 2/5
625/625 — 76s 122ms/step - accuracy: 0.5917 - loss: 0.6437 - val_accuracy: 0.5052 - val_loss: 0.7425
Epoch 3/5
625/625 — 83s 123ms/step - accuracy: 0.5773 - loss: 0.6641 - val_accuracy: 0.5048 - val_loss: 0.7193
```

RNN training time: 241.44 seconds

TRAINING LSTM


```
Epoch 1/5
625/625 — 178s 282ms/step - accuracy: 0.5332 - loss: 0.6862 - val_accuracy: 0.5406 - val_loss: 0.6786
Epoch 2/5
625/625 — 174s 279ms/step - accuracy: 0.6546 - loss: 0.6073 - val_accuracy: 0.7472 - val_loss: 0.5492
Epoch 3/5
625/625 — 176s 282ms/step - accuracy: 0.8042 - loss: 0.4780 - val_accuracy: 0.5122 - val_loss: 0.8245
Epoch 4/5
625/625 — 205s 286ms/step - accuracy: 0.6089 - loss: 0.6242 - val_accuracy: 0.7712 - val_loss: 0.5531
```


LSTM training time: 733.08 seconds

TRAINING GRU

```
Epoch 1/5
625/625 — 189s 298ms/step - accuracy: 0.5123 - loss: 0.6954 - val_accuracy: 0.5096 - val_loss: 0.6919
Epoch 2/5
625/625 — 186s 298ms/step - accuracy: 0.6488 - loss: 0.5958 - val_accuracy: 0.8454 - val_loss: 0.3722
Epoch 3/5
625/625 — 191s 306ms/step - accuracy: 0.8908 - loss: 0.2739 - val_accuracy: 0.8862 - val_loss: 0.2841
Epoch 4/5
625/625 — 187s 299ms/step - accuracy: 0.9471 - loss: 0.1552 - val_accuracy: 0.8864 - val_loss: 0.3115
```

Training Accuracy Comparison





```
results = []
predictions = {}
probabilities = {}

for name, model in models.items():

    print("\n")
    print("=" * 70)
    print(f"EVALUATING {name}")
    print("=" * 70)

    # Probability predictions
    y_prob = model.predict(
        X_test_pad,
        verbose=0
    ).ravel()

    # Convert probabilities to class labels
    y_pred = (y_prob >= 0.5).astype(int)

    probabilities[name] = y_prob
    predictions[name] = y_pred

    # Metrics
    accuracy = accuracy_score(
        y_test,
        y_pred
    )


    precision = precision_score(
        y_test,
        y_pred
    )

    recall = recall_score(
        y_test,
        y_pred
    )
```

EVALUATING GRU

	Model	Accuracy	Precision	Recall	F1 Score	Training Time	Parameters
0	RNN	0.5122	0.5157	0.4008	0.4510	241.4445	2572417
1	LSTM	0.7460	0.7079	0.8376	0.7673	733.0795	2609473
2	GRU	0.8814	0.8549	0.9188	0.8857	954.4171	2597313

RNN vs LSTM vs GRU — Performance Comparison



Test Accuracy Comparison

